

Session 06 Phase 3 – Master Results Tables

July 31, 2026

What each condition means

Every row below is a different way of turning the model’s internal activity while it writes an answer into a single feature vector. Think of that internal activity as a trail of numbers – one per layer, per generated token – and each condition as a different way of boiling that trail down.

core_max The baseline. Takes the single highest value each internal signal ever reaches while generating the answer – a snapshot of the loudest moment.

q_velocity Instead of the activity itself, tracks how much it *changes* from one layer to the next – speed instead of position. Picks up on whether the model’s internal representation is still being revised as information moves through the network, not just where it ends up.

q_static The same kind of peak snapshot as `core_max`, but softer: uses a high/low percentile instead of the literal maximum/minimum, so one unusual token can’t skew the whole signal.

core_concat Staples `core_max`’s peak snapshot and `q_velocity`’s change signal together into one longer feature vector – “where it ended up” plus “how it got there.”

joint_tensor Instead of computing two signals separately and combining them afterward, fits one shared compression across the combined raw data from the start, so the signals can influence each other during compression rather than only being stapled together at the end.

triple_concat Staples all three single signals – `core_max`, `q_static`, and `q_velocity` – together. The “use everything we have” condition.

Model	Condition	TruthfulQA	TriviaQA	NQ-Open	TyDiQA-GP
LLaMA-3.1-8B	HARP (published, supervised)	88.5	92.9	89.4	86.6
	core_max	84.4	86.4	85.7	83.1
	q_velocity	86.2	85.8	84.6	82.6
	q_static	84.3	86.7	86.1	84.0
	core_concat	86.0	86.8 [†]	85.3	83.2
	joint_tensor	84.5	86.4	86.3	83.3
	triple_concat	85.5	87.0	85.7	83.2
Qwen2.5-7B	HARP (published, supervised)	88.1	92.8	84.0	88.4
	core_max	87.1	91.3 [†]	78.8	88.3[‡]
	q_velocity	87.7	<i>pending</i>	78.2	87.0
	q_static	87.1	<i>pending</i>	79.1	87.5
	core_concat	87.8	<i>pending</i>	79.4	87.8
	joint_tensor	87.2	<i>pending</i>	78.5	87.0
	triple_concat	87.8	<i>pending</i>	79.3	87.9

Table 1: **Literature-comparable results: pooled AUROC (%) under HARP’s evaluation protocol** (train on 75% of known questions; test on the rest plus all fully-hallucinated questions), 5-seed mean, Random Forest readout. All non-gray rows are *our* unsupervised feature variants scored under this protocol; the gray row is the published score of HARP’s own supervised detector, copied from their paper for reference (not recomputed by us). Bold = best of our conditions per column (HARP row excluded). [†]On TriviaQA – and *only* on TriviaQA – logistic regression outperforms the forest, for *both* models, which makes this a property of the dataset rather than the architecture. LLaMA: core_max 87.3 ± 0.2 , q_static 87.3 ± 0.3 , core_concat 87.3 ± 0.2 . Qwen: core_max 92.4 ± 0.1 , our single best literature-comparable number in the study and 0.4 points below HARP’s published 92.8. [‡]Qwen/TyDiQA-GP is the first cell where an unsupervised condition of ours reaches the published supervised score: core_max 88.3 ± 1.8 against HARP’s reported 88.4, a gap of 0.1 points, well inside our seed-to-seed spread.

Model	Condition	TruthfulQA	TriviaQA	NQ-Open	TyDiQA-GP
LLaMA-3.1-8B	core_max	81.77	81.97	76.10	84.48
	q_velocity	83.58	81.17	75.44	85.68
	q_static	83.28	81.97	76.72	85.26
	core_concat	84.35*	82.06	76.74	85.27
	joint_tensor	83.76	81.84	75.70	85.13
	triple_concat	83.28	82.27	76.01	85.76
Qwen2.5-7B	core_max	82.47	88.56	74.45	87.39
	q_velocity	84.33	<i>pending</i>	74.88	88.50
	q_static	83.28	<i>pending</i>	75.32	86.88
	core_concat	84.43	<i>pending</i>	76.18	88.85
	joint_tensor	82.69	<i>pending</i>	75.87	87.11
	triple_concat	84.05	<i>pending</i>	77.61	89.14

Table 2: **Within-prompt AUROC (%)**, grouped 5-fold CV, **Random Forest** – our stricter internal metric: among the 10 answers to the *same* question, can the detector rank the hallucinated ones above the truthful ones? No HARP reference row appears here because this metric is not reported in their paper. Bold = best per column, but note that a bolded lead is *nominal* unless marked *: on TriviaQA, NQ-Open, and TyDiQA-GP no condition’s within-prompt gain over core_max survives a paired bootstrap (all 95% CIs include zero; triple_concat’s TriviaQA lead, for instance, is $\Delta = +0.30$ pts, CI $[-0.14, +0.73]$). *core_concat vs. core_max on TruthfulQA remains, across all four completed datasets, the project’s *only* statistically significant within-prompt improvement (paired $\Delta = +2.57$ pts, 95% CI $[+0.03, +5.26]$, excludes zero). Separately, all four combination conditions *do* significantly beat q_velocity on TriviaQA on both metrics.

Model	Condition	TruthfulQA		TriviaQA		NQ-Open		TyDiQA-GP	
		RF	LR	RF	LR	RF	LR	RF	LR
LLaMA-3.1-8B	core_max	✓			✓	✓		✓	
	q_velocity	✓		✓		✓		✓	
	q_static	✓			✓	✓		✓	
	core_concat	✓			✓	✓		✓	
	joint_tensor	✓			✓	✓		✓	
	triple_concat	✓			✓	✓		✓	
Qwen2.5-7B	core_max	✓			✓	✓		✓	
	q_velocity	✓			<i>pend</i>	✓		✓	
	q_static	✓			<i>pend</i>	✓		✓	
	core_concat	✓			<i>pend</i>	✓		✓	
	joint_tensor	✓			<i>pend</i>	✓		✓	
	triple_concat				<i>pending</i>			✓	

Table 3: **Which classifier wins, under the HARP protocol** – ✓ marks the higher of RF/LR (5-seed mean pooled AUROC) for that condition×dataset cell. RF sweeps TruthfulQA, NQ-Open, and TyDiQA-GP (18/18 LLaMA cells). TriviaQA flips: LR wins 5 of 6 conditions, RF holds only on q_velocity. See Table 1’s caption for the actual TriviaQA numbers (LR’s wins there are narrow, 0.2–0.9pt – not a landslide like RF’s wins elsewhere). Qwen reproduces this exactly: RF sweeps TruthfulQA, NQ-Open and TyDiQA-GP (18/18 cells), and its one completed TriviaQA condition flips to LR (92.4 vs 91.3). The RF/LR inversion therefore tracks the *dataset*, not the model – the same dataset flips both architectures.


Reading the two metrics

Both metrics answer “can the detector rank hallucinated answers above truthful ones,” but over different comparison sets, and the difference matters for how the numbers should be read.

- **Pooled AUROC.** Put every answer from every question into one pile and rank the whole pile. The catch: some questions are easy (the model gets all 10 tries right) and some are hard (it gets all 10 wrong). A detector can score well here just by recognising “this looks like a hard question,” without ever telling a good answer from a bad one *for the same question*. This is the number the literature reports, so we report it for comparability.
- **Within-prompt AUROC.** For each question separately, compare only that question’s own answers against each other, then average across questions. Questions where all 10 answers are right, or all 10 are wrong, contribute nothing – there is no correct/incorrect pair to rank. This removes the easy-vs-hard shortcut entirely, so it is the harder and more honest number, and it is what we select on.
- **“OOF” = out-of-fold.** With 5-fold cross-validation, every answer is scored by the one model that did *not* see it during training. “Pooled OOF AUROC” means: collect all those held-out scores, put them in one pile, and compute a single AUROC. It is not an average of five per-fold numbers – it is one number over all held-out predictions, so no answer is ever scored by a model that trained on it.

Figure 3 exists precisely to measure how much of the pooled number is the shortcut and how much is real detection.

Figure 1 – where in the network the signal lives




LLaMA-3.1-8B. Left: how much larger the average size of the internal activity is for hallucinated answers than for truthful ones, in pooled standard deviations. Right: detection score when the detector is allowed to use only that one layer (64 components, logistic regression). Bands are 95% confidence intervals. Only layers 15–23 exist on disk.

What it shows

- **Left panel: the curves are flat.** There is no bump at any layer. The “layer-22 hump” we have been assuming does not exist – the whole range of movement across nine layers is about 0.1 standard deviations, which is nothing.
- **The datasets sit far apart, and one has the opposite sign.** NQ-Open is at +1.3, TriviaQA +0.7, TyDiQA-GP +0.35, and TruthfulQA at −0.85 – there, hallucinated answers have *smaller* internal activity, not larger. So this quantity behaves like a fingerprint of the dataset, not a signal of hallucination.
- **Critically: this quantity does not predict detection at all.** NQ-Open has the largest gap and TruthfulQA has a negative one, yet both sit at the top of the right panel. Since the 15–23 window was originally chosen using exactly this kind of energy audit, the method we used to pick our window does not track what we actually care about.
- **Right panel: layer 15 is the best single layer on all four datasets.** Three of the four decline steadily to the right; TruthfulQA is flat rather than falling.
- **Layer 15 is the left edge of our window, which is a warning sign.** When the best value sits on the boundary of the range you looked at, the true best is usually outside it. The honest reading is that we may have been extracting from the wrong part of the network, and we cannot tell because the pipeline crops at 15. Settling this needs one GPU pass over earlier layers.
- **A single layer appears to match the entire pipeline.** Layer 15 alone, with 64 numbers and a simple linear model, scores about 0.877 / 0.862 / 0.873 / 0.830 across the four datasets, against 0.869 / 0.854 / 0.865 / 0.821 for the full 9-layer, 320-number random forest. The confidence intervals overlap, so the claim is “matches,” not “beats” – but even a match means the other eight layers are contributing close to nothing. This needs one cheap paired check before we state it publicly.

Figure 2 – do the static and velocity streams carry different things?




Left: for each pair of directions (1st most important, 2nd, and so on up to 64th), how closely aligned the static and velocity subspaces are. 1.0 means the two directions are identical; 0 means completely unrelated. Right: how stable the static subspace is, comparing folds within a dataset (diagonal) and datasets against each other (off-diagonal).

What it shows

- **Three clearly separated bands in the left panel.** The black line (the two halves of the *same* stream, q95 vs q05) stays near 1.0 almost the whole way – they are near-copies, sharing about 58 of 64 directions. The coloured lines (static vs velocity) drop below 0.5 around the 15th direction – about 15 of 64 shared. The grey line (two random subspaces) never rises above 0.18 – nothing shared.
- **So static and velocity really are geometrically different** – far more different than chance, far less different than two halves of one stream. That part of the story holds.
- **But being different did not make them useful together.** Combining the two streams buys only 0.1–2.5 points, mostly not statistically significant. So the two streams point in different directions while carrying largely the *same* information about hallucination. This is the honest headline, and it explains why several rounds of feature engineering produced about one point.
- **Where the overlap sits matters.** The first five directions are nearly identical (above 0.8). The part the two streams share is their biggest, most dominant part – which is plausibly the part a classifier leans on most. The genuinely different directions are the small, tail ones.
- **A free simplification.** Because q05 is nearly a copy of q95, we could likely drop q05 entirely and halve the feature size at no cost. That is a CPU-only check on data we already have.
- **Right panel, and a good rigour point:** comparing two different folds of the same dataset gives 0.99–1.00. Our unsupervised basis is essentially deterministic – it does not wobble when the data is resampled. Across different datasets it is 0.69–0.92, so there is substantial shared structure, with TyDiQA-GP the most distinct.
- **One caution when presenting:** the diagonal and off-diagonal cells of that heatmap measure two different comparisons. It is stated in the title but is easy to miss.

Figure 3 – what the standard metric is actually measuring



Left: pooled AUROC using full features, using only each question’s 10-answer average (“centroid-only,” which destroys all per-answer information), and using each answer minus that average (“delta-only,” which destroys all question-identity information). Right: pooled AUROC minus within-prompt AUROC, per dataset.

What it shows

- **Question identity alone gets most of the way there.** The centroid-only bars reach 0.71–0.81. Those features contain *zero* information about any individual answer – every answer to a question has been replaced by the same average. Yet they recover 66–85% of the pooled score’s lift above chance.
- **What this does and does not mean.** It does not mean the metric is invalid – those features still come from the model’s own internal states. It means pooled AUROC is mostly measuring *which questions the model will fail on*, not *which specific answers are lies*. Most readers of a hallucination-detection paper assume the second.
- **The size of the distortion varies a lot by benchmark** (right panel). NQ-Open is the most inflated at +0.108, then TruthfulQA +0.049 and TriviaQA +0.034. TyDiQA-GP is the only one where pooled *understates* performance, at −0.025.
- **The most useful result is not currently drawn in the figure.** Removing the question average and keeping only the per-answer part *beats* the full features on within-prompt AUROC: 83.2 vs 82.0 on TruthfulQA, 83.8 vs 82.0 on TriviaQA. So the question average is not just unhelpful for per-answer detection – it is actively getting in the way.
- **That is a plain subtraction beating all six engineered feature variants** on the metric we care about, at no compute cost. It should go into the condition grid for both models.
- **One honest caveat on it.** Subtracting the average needs the other nine answers available at test time, so it is a “multi-sample” method – the same setting as SelfCheckGPT and INSIDE, which is legitimate and precedented, but it is not a like-for-like comparison against single-pass HARP and must be labelled as such. It uses no labels, so there is no leakage.
- **This lines up with the tables.** Combination conditions beat the baseline on pooled but never on within-prompt, for both models and every dataset – exactly what you would expect if those combinations are adding question-difficulty information rather than per-answer signal. Two separate analyses reaching the same conclusion.